

# Kaggle API setup

In [3]: `!pip install kaggle`

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.7/site-packages (1.5.12)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.7/site-packages (from kaggle) (1.15.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/site-packages (from kaggle) (4.62.1)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.7/site-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.7/site-packages (from kaggle) (2.25.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.7/site-packages (from kaggle) (5.0.2)
Requirement already satisfied: certifi in /usr/local/lib/python3.7/site-packages (from kaggle) (2021.5.30)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/site-packages (from kaggle) (1.26.6)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.7/site-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/site-packages (from requests->kaggle) (2.10)
Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.7/site-packages (from requests->kaggle) (4.0.0)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
WARNING: You are using pip version 21.2.4; however, version 21.3.1 is available.
You should consider upgrading via the '/usr/local/bin/python -m pip install --upgrade pip' command.
```

In [19]: `!pip install pandoc`

```
Collecting pandoc
  Downloading pandoc-1.1.0-py3-none-any.whl (27 kB)
Collecting ply
  Downloading ply-3.11-py2.py3-none-any.whl (49 kB)
    |████████████████████████████████████████| 49 kB 2.2 MB/s eta 0:00:01
Collecting plumbum
  Downloading plumbum-1.7.0-py2.py3-none-any.whl (116 kB)
    |████████████████████████████████████████| 116 kB 4.9 MB/s eta 0:00:01
Installing collected packages: ply, plumbum, pandoc
Successfully installed pandoc-1.1.0 plumbum-1.7.0 ply-3.11
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
WARNING: You are using pip version 21.2.4; however, version 21.3.1 is available.
You should consider upgrading via the '/usr/local/bin/python -m pip install --upgrade pip' command.
```

In [4]: `!pwd`

```
/root/shared/Courses/AML526/I526_AML_Student/Assignments/Unit-Project-Home-Credit-Default-Risk/HCDR_Phase1_baseline_submission
```



```

import os
import zipfile
from sklearn.base import BaseEstimator, TransformerMixin
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler
from sklearn.pipeline import Pipeline, FeatureUnion
from pandas.plotting import scatter_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
import os
import zipfile
from sklearn.base import BaseEstimator, TransformerMixin
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler
from sklearn.pipeline import Pipeline, FeatureUnion
from pandas.plotting import scatter_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, ShuffleSplit, cross_val_score
from sklearn.metrics import roc_curve, roc_auc_score, accuracy_score
from scipy import stats
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')

```

## Data Loads

### Application\_Train data load

In [5]:

```

def load_data(in_path, name):
    df = pd.read_csv(in_path)
    print(f"{name}: shape is {df.shape}")
    print(df.info())
    display(df.head(5))
    return df

datasets={} # Lets store the datasets in a dictionary so we can keep track of them eas

```


```
ds_name = 'application_train'
datasets[ds_name] = load_data(os.path.join(DATA_DIR, f'{ds_name}.csv'), ds_name)

datasets['application_train'].shape
```

```
application_train: shape is (307511, 122)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 307511 entries, 0 to 307510
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(41), object(16)
memory usage: 286.2+ MB
None
```

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY
0	100002	1	Cash loans	M	N	Y
1	100003	0	Cash loans	F	N	N
2	100004	0	Revolving loans	M	Y	Y
3	100006	0	Cash loans	F	N	Y
4	100007	0	Cash loans	M	N	Y

5 rows × 122 columns



Out[5]: (307511, 122)


## Application\_Test data load

```
In [7]: ds_name = 'application_test'
datasets[ds_name] = load_data(os.path.join(DATA_DIR, f'{ds_name}.csv'), ds_name)
```

```
application_test: shape is (48744, 121)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48744 entries, 0 to 48743
Columns: 121 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(40), object(16)
memory usage: 45.0+ MB
None
```

	SK_ID_CURR	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHI
0	100001	Cash loans	F	N	Y	
1	100005	Cash loans	M	N	Y	
2	100013	Cash loans	M	Y	Y	
3	100028	Cash loans	F	N	Y	
4	100038	Cash loans	M	Y	N	

5 rows × 121 columns



## Other Datasets load

In [8]:

```
%%time
ds_names = ("application_train", "application_test", "bureau", "bureau_balance", "credit_
            "previous_application", "POS_CASH_balance")

for ds_name in ds_names:
    datasets[ds_name] = load_data(os.path.join(DATA_DIR, f'{ds_name}.csv'), ds_name)
```

application\_train: shape is (307511, 122)  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 307511 entries, 0 to 307510  
Columns: 122 entries, SK\_ID\_CURR to AMT\_REQ\_CREDIT\_BUREAU\_YEAR  
dtypes: float64(65), int64(41), object(16)  
memory usage: 286.2+ MB  
None

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY
0	100002	1	Cash loans	M	N	Y
1	100003	0	Cash loans	F	N	N
2	100004	0	Revolving loans	M	Y	Y
3	100006	0	Cash loans	F	N	Y
4	100007	0	Cash loans	M	N	Y

5 rows × 122 columns

```
application_test: shape is (48744, 121)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48744 entries, 0 to 48743
Columns: 121 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(40), object(16)
memory usage: 45.0+ MB
None
```

	SK_ID_CURR	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHI
0	100001	Cash loans	F	N	Y	
1	100005	Cash loans	M	N	Y	
2	100013	Cash loans	M	Y	Y	
3	100028	Cash loans	F	N	Y	
4	100038	Cash loans	M	Y	N	

5 rows × 121 columns

```
bureau: shape is (1716428, 17)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1716428 entries, 0 to 1716427
Data columns (total 17 columns):
#   Column              Dtype
---  -
0   SK_ID_CURR          int64
1   SK_ID_BUREAU        int64
2   CREDIT_ACTIVE       object
```

```
3 CREDIT_CURRENCY      object
4 DAYS_CREDIT          int64
5 CREDIT_DAY_OVERDUE   int64
6 DAYS_CREDIT_ENDDATE  float64
7 DAYS_ENDDATE_FACT    float64
8 AMT_CREDIT_MAX_OVERDUE float64
9 CNT_CREDIT_PROLONG   int64
10 AMT_CREDIT_SUM       float64
11 AMT_CREDIT_SUM_DEBT  float64
12 AMT_CREDIT_SUM_LIMIT float64
13 AMT_CREDIT_SUM_OVERDUE float64
14 CREDIT_TYPE          object
15 DAYS_CREDIT_UPDATE   int64
16 AMT_ANNUITY          float64
dtypes: float64(8), int64(6), object(3)
memory usage: 222.6+ MB
None
```

	SK_ID_CURR	SK_ID_BUREAU	CREDIT_ACTIVE	CREDIT_CURRENCY	DAYS_CREDIT	CREDIT_DAY_OVERDUE
0	215354	5714462	Closed	currency 1	-497	
1	215354	5714463	Active	currency 1	-208	
2	215354	5714464	Active	currency 1	-203	
3	215354	5714465	Active	currency 1	-203	
4	215354	5714466	Active	currency 1	-629	

```
bureau_balance: shape is (27299925, 3)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27299925 entries, 0 to 27299924
Data columns (total 3 columns):
#   Column      Dtype
---  -
0   SK_ID_BUREAU  int64
1   MONTHS_BALANCE int64
2   STATUS       object
dtypes: int64(2), object(1)
memory usage: 624.8+ MB
None
```

	SK_ID_BUREAU	MONTHS_BALANCE	STATUS
0	5715448	0	C
1	5715448	-1	C
2	5715448	-2	C
3	5715448	-3	C
4	5715448	-4	C

```
credit_card_balance: shape is (3840312, 23)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3840312 entries, 0 to 3840311
Data columns (total 23 columns):
#   Column      Dtype
```

```

---  -----
0  SK_ID_PREV          int64
1  SK_ID_CURR          int64
2  MONTHS_BALANCE      int64
3  AMT_BALANCE          float64
4  AMT_CREDIT_LIMIT_ACTUAL int64
5  AMT_DRAWINGS_ATM_CURRENT float64
6  AMT_DRAWINGS_CURRENT float64
7  AMT_DRAWINGS_OTHER_CURRENT float64
8  AMT_DRAWINGS_POS_CURRENT float64
9  AMT_INST_MIN_REGULARITY float64
10 AMT_PAYMENT_CURRENT float64
11 AMT_PAYMENT_TOTAL_CURRENT float64
12 AMT_RECEIVABLE_PRINCIPAL float64
13 AMT_RECIVABLE        float64
14 AMT_TOTAL_RECEIVABLE float64
15 CNT_DRAWINGS_ATM_CURRENT float64
16 CNT_DRAWINGS_CURRENT int64
17 CNT_DRAWINGS_OTHER_CURRENT float64
18 CNT_DRAWINGS_POS_CURRENT float64
19 CNT_INSTALMENT_MATURE_CUM float64
20 NAME_CONTRACT_STATUS object
21 SK_DPD               int64
22 SK_DPD_DEF           int64
dtypes: float64(15), int64(7), object(1)
memory usage: 673.9+ MB
None
```

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	AMT_BALANCE	AMT_CREDIT_LIMIT_ACTUAL	AMT_DR
0	2562384	378907	-6	56.970	135000	
1	2582071	363914	-1	63975.555	45000	
2	1740877	371185	-7	31815.225	450000	
3	1389973	337855	-4	236572.110	225000	
4	1891521	126868	-1	453919.455	450000	

5 rows × 23 columns

```

installments_payments: shape is (13605401, 8)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13605401 entries, 0 to 13605400
Data columns (total 8 columns):
#   Column                                Dtype
---  -----
0   SK_ID_PREV                            int64
1   SK_ID_CURR                            int64
2   NUM_INSTALMENT_VERSION                float64
3   NUM_INSTALMENT_NUMBER                 int64
4   DAYS_INSTALMENT                       float64
5   DAYS_ENTRY_PAYMENT                    float64
6   AMT_INSTALMENT                        float64
7   AMT_PAYMENT                           float64
dtypes: float64(5), int64(3)
memory usage: 830.4 MB
None
```

SK_ID_PREV	SK_ID_CURR	NUM_INSTALMENT_VERSION	NUM_INSTALMENT_NUMBER	DAYS_INSTALMI
------------	------------	------------------------	-----------------------	---------------

	SK_ID_PREV	SK_ID_CURR	NUM_INSTALMENT_VERSION	NUM_INSTALMENT_NUMBER	DAYS_INSTALMENT
0	1054186	161674	1.0	6	-118
1	1330831	151639	0.0	34	-215
2	2085231	193053	2.0	1	-6
3	2452527	199697	1.0	3	-24
4	2714724	167756	1.0	2	-138

previous\_application: shape is (1670214, 37)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1670214 entries, 0 to 1670213

Data columns (total 37 columns):

#	Column	Non-Null Count	Dtype
0	SK_ID_PREV	1670214 non-null	int64
1	SK_ID_CURR	1670214 non-null	int64
2	NAME_CONTRACT_TYPE	1670214 non-null	object
3	AMT_ANNUITY	1297979 non-null	float64
4	AMT_APPLICATION	1670214 non-null	float64
5	AMT_CREDIT	1670213 non-null	float64
6	AMT_DOWN_PAYMENT	774370 non-null	float64
7	AMT_GOODS_PRICE	1284699 non-null	float64
8	WEEKDAY_APPR_PROCESS_START	1670214 non-null	object
9	HOURLY_APPR_PROCESS_START	1670214 non-null	int64
10	FLAG_LAST_APPL_PER_CONTRACT	1670214 non-null	object
11	NFLAG_LAST_APPL_IN_DAY	1670214 non-null	int64
12	RATE_DOWN_PAYMENT	774370 non-null	float64
13	RATE_INTEREST_PRIMARY	5951 non-null	float64
14	RATE_INTEREST_PRIVILEGED	5951 non-null	float64
15	NAME_CASH_LOAN_PURPOSE	1670214 non-null	object
16	NAME_CONTRACT_STATUS	1670214 non-null	object
17	DAYS_DECISION	1670214 non-null	int64
18	NAME_PAYMENT_TYPE	1670214 non-null	object
19	CODE_REJECT_REASON	1670214 non-null	object
20	NAME_TYPE_SUITE	849809 non-null	object
21	NAME_CLIENT_TYPE	1670214 non-null	object
22	NAME_GOODS_CATEGORY	1670214 non-null	object
23	NAME_PORTFOLIO	1670214 non-null	object
24	NAME_PRODUCT_TYPE	1670214 non-null	object
25	CHANNEL_TYPE	1670214 non-null	object
26	SELLERPLACE_AREA	1670214 non-null	int64
27	NAME_SELLER_INDUSTRY	1670214 non-null	object
28	CNT_PAYMENT	1297984 non-null	float64
29	NAME_YIELD_GROUP	1670214 non-null	object
30	PRODUCT_COMBINATION	1669868 non-null	object
31	DAYS_FIRST_DRAWING	997149 non-null	float64
32	DAYS_FIRST_DUE	997149 non-null	float64
33	DAYS_LAST_DUE_1ST_VERSION	997149 non-null	float64
34	DAYS_LAST_DUE	997149 non-null	float64
35	DAYS_TERMINATION	997149 non-null	float64
36	NFLAG_INSURED_ON_APPROVAL	997149 non-null	float64

dtypes: float64(15), int64(6), object(16)

memory usage: 471.5+ MB

None

SK_ID_PREV	SK_ID_CURR	NAME_CONTRACT_TYPE	AMT_ANNUITY	AMT_APPLICATION	AMT_CREDIT
------------	------------	--------------------	-------------	-----------------	------------



	SK_ID_PREV	SK_ID_CURR	NAME_CONTRACT_TYPE	AMT_ANNUITY	AMT_APPLICATION	AMT_CREDIT
0	2030495	271877	Consumer loans	1730.430	17145.0	17145.0
1	2802425	108129	Cash loans	25188.615	607500.0	679671.0
2	2523466	122040	Cash loans	15060.735	112500.0	136444.5
3	2819243	176158	Cash loans	47041.335	450000.0	470790.0
4	1784265	202054	Cash loans	31924.395	337500.0	404055.0

5 rows × 37 columns

POS\_CASH\_balance: shape is (10001358, 8)  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10001358 entries, 0 to 10001357  
Data columns (total 8 columns):  
# Column Dtype  
--- -  
0 SK\_ID\_PREV int64  
1 SK\_ID\_CURR int64  
2 MONTHS\_BALANCE int64  
3 CNT\_INSTALMENT float64  
4 CNT\_INSTALMENT\_FUTURE float64  
5 NAME\_CONTRACT\_STATUS object  
6 SK\_DPD int64  
7 SK\_DPD\_DEF int64  
dtypes: float64(2), int64(5), object(1)  
memory usage: 610.4+ MB  
None

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	CNT_INSTALMENT	CNT_INSTALMENT_FUTURE	NAME_
0	1803195	182943	-31	48.0	45.0	
1	1715348	367990	-33	36.0	35.0	
2	1784872	397406	-32	12.0	9.0	
3	1903291	269225	-35	48.0	42.0	
4	2341044	334279	-35	36.0	35.0	

CPU times: user 45.9 s, sys: 7.47 s, total: 53.4 s  
Wall time: 1min 30s

# EDA

## Missing data for application train

In [9]:

```
percent = (datasets["application_train"].isnull().sum()/datasets["application_train"].i
sum_missing = datasets["application_train"].isna().sum().sort_values(ascending = False)
missing_application_train_data = pd.concat([percent, sum_missing], axis=1, keys=['Perc
missing_application_train_data.head(20)
```

Out[9]:

Percent	Train Missing Count
---------	---------------------

	Percent	Train Missing Count
<b>COMMONAREA_MEDI</b>	69.87	214865
<b>COMMONAREA_AVG</b>	69.87	214865
<b>COMMONAREA_MODE</b>	69.87	214865
<b>NONLIVINGAPARTMENTS_MODE</b>	69.43	213514
<b>NONLIVINGAPARTMENTS_AVG</b>	69.43	213514
<b>NONLIVINGAPARTMENTS_MEDI</b>	69.43	213514
<b>FONDKAPREMONT_MODE</b>	68.39	210295
<b>LIVINGAPARTMENTS_MODE</b>	68.35	210199
<b>LIVINGAPARTMENTS_AVG</b>	68.35	210199
<b>LIVINGAPARTMENTS_MEDI</b>	68.35	210199
<b>FLOORSMIN_AVG</b>	67.85	208642
<b>FLOORSMIN_MODE</b>	67.85	208642
<b>FLOORSMIN_MEDI</b>	67.85	208642
<b>YEARS_BUILD_MEDI</b>	66.50	204488
<b>YEARS_BUILD_MODE</b>	66.50	204488
<b>YEARS_BUILD_AVG</b>	66.50	204488
<b>OWN_CAR_AGE</b>	65.99	202929
<b>LANDAREA_MEDI</b>	59.38	182590
<b>LANDAREA_MODE</b>	59.38	182590
<b>LANDAREA_AVG</b>	59.38	182590

## Missing data for application test

```
In [10]: percent = (datasets["application_test"].isnull().sum()/datasets["application_test"].isn
sum_missing = datasets["application_test"].isna().sum().sort_values(ascending = False)
missing_application_train_data = pd.concat([percent, sum_missing], axis=1, keys=['Perc
missing_application_train_data.head(20)
```


```
Out[10]:
```

	Percent	Test Missing Count
<b>COMMONAREA_AVG</b>	68.72	33495
<b>COMMONAREA_MODE</b>	68.72	33495
<b>COMMONAREA_MEDI</b>	68.72	33495
<b>NONLIVINGAPARTMENTS_AVG</b>	68.41	33347
<b>NONLIVINGAPARTMENTS_MODE</b>	68.41	33347
<b>NONLIVINGAPARTMENTS_MEDI</b>	68.41	33347

	Percent	Test Missing Count
<b>FONDKAPREMONT_MODE</b>	67.28	32797
<b>LIVINGAPARTMENTS_AVG</b>	67.25	32780
<b>LIVINGAPARTMENTS_MODE</b>	67.25	32780
<b>LIVINGAPARTMENTS_MEDI</b>	67.25	32780
<b>FLOORSMIN_MEDI</b>	66.61	32466
<b>FLOORSMIN_AVG</b>	66.61	32466
<b>FLOORSMIN_MODE</b>	66.61	32466
<b>OWN_CAR_AGE</b>	66.29	32312
<b>YEARS_BUILD_AVG</b>	65.28	31818
<b>YEARS_BUILD_MEDI</b>	65.28	31818
<b>YEARS_BUILD_MODE</b>	65.28	31818
<b>LANDAREA_MEDI</b>	57.96	28254
<b>LANDAREA_AVG</b>	57.96	28254
<b>LANDAREA_MODE</b>	57.96	28254

## Distribution of the target column

In [23]: `datasets["application_train"]['TARGET'].astype(int).plot.hist();`



## Correlation with the target column

In [24]: `correlations = datasets["application_train"].corr()['TARGET'].sort_values()
print('Most Positive Correlations:\n', correlations.tail(10))
print('\nMost Negative Correlations:\n', correlations.head(10))`

```
Most Positive Correlations:
FLAG_DOCUMENT_3          0.044346
REG_CITY_NOT_LIVE_CITY   0.044395
```

```

FLAG_EMP_PHONE      0.045982
REG_CITY_NOT_WORK_CITY 0.050994
DAYS_ID_PUBLISH     0.051457
DAYS_LAST_PHONE_CHANGE 0.055218
REGION_RATING_CLIENT 0.058899
REGION_RATING_CLIENT_W_CITY 0.060893
DAYS_BIRTH          0.078239
TARGET              1.000000

```

Name: TARGET, dtype: float64

Most Negative Correlations:

```

EXT_SOURCE_3      -0.178919
EXT_SOURCE_2      -0.160472
EXT_SOURCE_1      -0.155317
DAYS_EMPLOYED     -0.044932
FLOORSMAX_AVG     -0.044003
FLOORSMAX_MEDI    -0.043768
FLOORSMAX_MODE    -0.043226
AMT_GOODS_PRICE   -0.039645
REGION_POPULATION_RELATIVE -0.037227
ELEVATORS_AVG     -0.034199

```

Name: TARGET, dtype: float64

```


In [12]: app_train=datasets["application_train"]
         app_test=datasets["application_test"]

```

```

In [26]: sns.countplot(data = app_train, x ='CODE_GENDER')
         plt.title('Distribution of Gender')
         plt.show()

```




It looks like we have more female Applicants when compared to male

```

In [29]: app_train['AGE'] = app_train['DAYS_BIRTH']/-365
         fig = sns.FacetGrid(app_train, col='TARGET', hue='TARGET', height=4)
         fig.map(sns.histplot, 'AGE', bins=30, kde=False)
         plt.show()

```



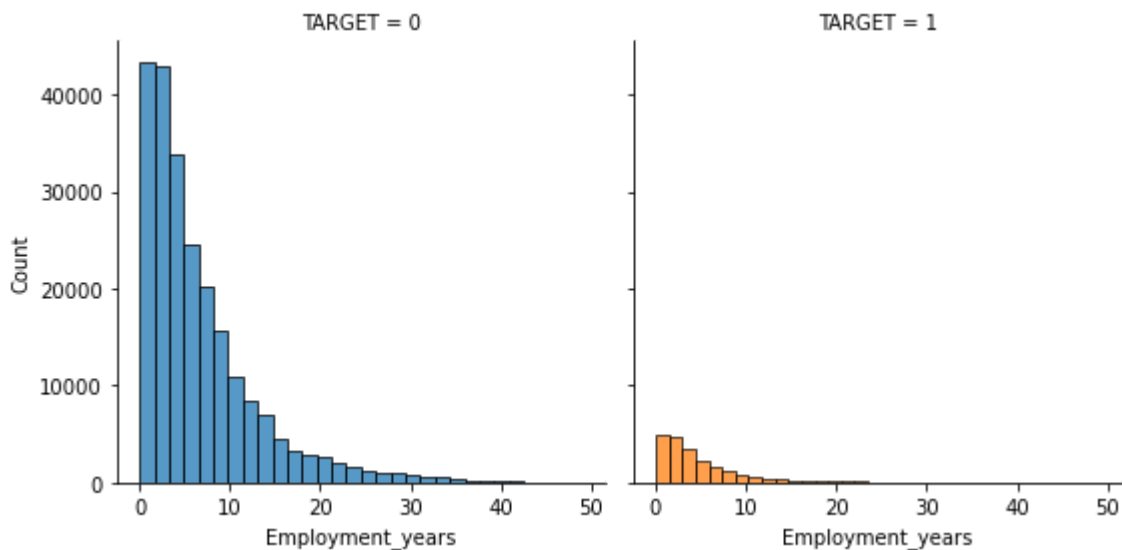
It can be observed that most of the defaulters are approximately 27-40 yrs old.

```
In [32]: # replace the incorrect values of Days_Employed
import numpy as np

# Create an error flag column
app_train['DAYS_EMPLOYED_ERROR'] = app_train["DAYS_EMPLOYED"] == 365243
app_test['DAYS_EMPLOYED_ERROR'] = app_test["DAYS_EMPLOYED"] == 365243 # do the same for

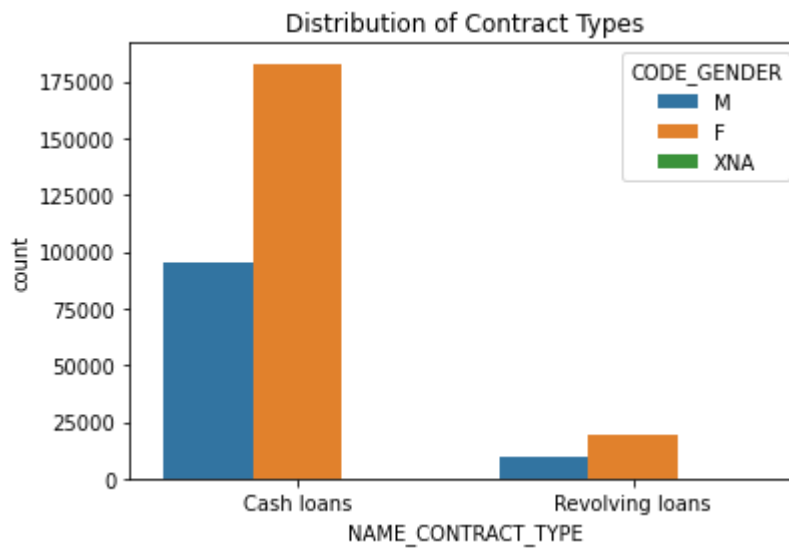
# Replace the error values with nan
app_train['DAYS_EMPLOYED'].replace({365243: np.nan}, inplace = True)

app_train['Employment_years'] = app_train['DAYS_EMPLOYED']/365
fig = sns.FacetGrid(app_train, col='TARGET', hue='TARGET', height=4)
fig.map(sns.histplot, 'Employment_years', bins=30, kde=False)
plt.show()
```



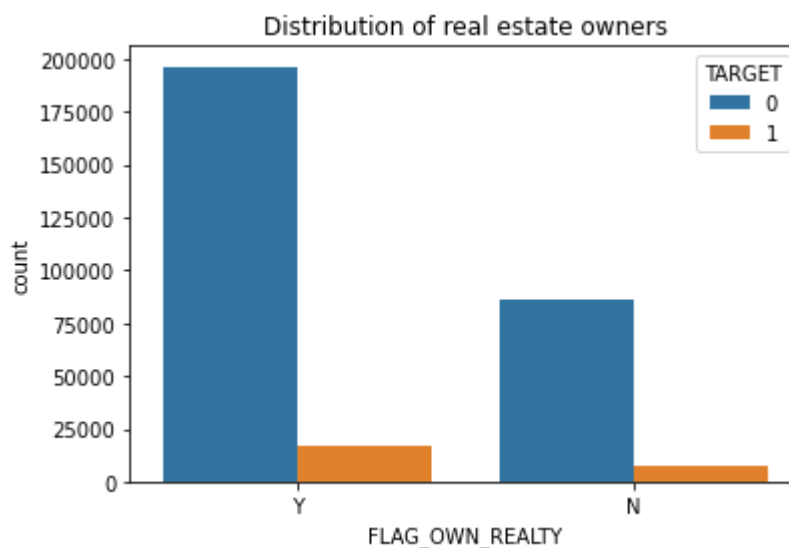
Applicants with less than 2 years of employment are less likely to repay the loan.

```
In [34]: sns.countplot(data = app_train, x = 'NAME_CONTRACT_TYPE', hue = 'CODE_GENDER')
plt.title('Distribution of Contract Types')
plt.show()
```




It can be seen that in general people mostly go for Cash loans as compared to revolving loans and both type of contracts are dominated by females.

```
In [36]: sns.countplot(data = app_train, x = 'FLAG_OWN_REALTY', hue = 'TARGET')
plt.title('Distribution of real estate owners')
plt.show()
```



Generally people who own a realty are more likely to go for loans as compared to people who don't own one.


```
In [13]: sns.countplot(data = app_train, x = 'NAME_INCOME_TYPE', hue = 'TARGET')
plt.title('Distribution of Income Sources ')
plt.xticks(rotation = 90)
plt.show()
```



It can be observed that most of the defaulters belong to the working class followed by Commercial associates. A state servant is the least likely to fail repaying the loan amount

In [40]:


```
sns.countplot(data = app_train, x = 'OCCUPATION_TYPE', hue = 'TARGET')
plt.title('Occupation types ')
plt.xticks(rotation = 90)
plt.show()
```



People who do not provide their occupation details and those who are labourers are more likely to fail paying back the loan.

In [42]:

```
ig, ax = plt.subplots(figsize=(15, 6))
ax = sns.boxplot(y = "OCCUPATION_TYPE", x = "AMT_INCOME_TOTAL", orient = "h", data = ap)
plt.xlim([0, 1e6])
plt.show()
```



## Processing Pipeline

In [64]:

```
class DataFrameSelector(BaseEstimator, TransformerMixin):
    def __init__(self, attribute_names):
        self.attribute_names = attribute_names
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        return X[self.attribute_names].values
```

In [71]:

```
def previous_app():
    df_pa = datasets['previous_application']
    #df_pa.drop('SK_ID_PREV', axis = 1, inplace = True)

    numcols = df_pa.select_dtypes(exclude= 'object').columns
    catcols = df_pa.select_dtypes(include= 'object').columns

    num = df_pa[numcols]
    cat = df_pa[catcols]

    num.index = num['SK_ID_CURR']
    cat.index = num['SK_ID_CURR']

    catind = cat.index

    num.drop('SK_ID_CURR', inplace= True, axis = 1)
    num = num.groupby('SK_ID_CURR').agg([np.sum, np.mean, max, min])
    num.columns = num.columns.map('_'.join)
    mdl = StandardScaler().fit(num)
    num = pd.DataFrame(mdl.transform(num), index = num.index, columns = num.columns)
```



```

cat = SimpleImputer(strategy='constant').fit_transform(cat)
mdl = OneHotEncoder().fit(cat)

cat = pd.DataFrame(mdl.transform(cat).toarray(), index = catind, columns = mdl.get_
cat = cat.groupby('SK_ID_CURR').agg(np.mean)
final = num.merge(cat, how = 'left', on = 'SK_ID_CURR')
return final

```

In [66]:

```

def load_train():
    df_train = app_train

    selected_features = ['SK_ID_CURR', 'AMT_INCOME_TOTAL', 'AMT_CREDIT', 'DAYS_EMPLOYED',
                        'EXT_SOURCE_2', 'EXT_SOURCE_3', 'CODE_GENDER', 'FLAG_OWN_REALTY', 'FLA
                        'NAME_CONTRACT_TYPE', 'NAME_EDUCATION_TYPE', 'OCCUPATION_TYPE', 'NAME

    df_train['DAYS_EMPLOYED'].replace({365243: np.nan}, inplace = True)
    df_train['DAYS_BIRTH'] = df_train['DAYS_BIRTH']/-365

    X_train= df_train[selected_features]
    y_train = df_train['TARGET']

    X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train, test_size=0

    X_val_ind = X_valid['SK_ID_CURR']

    X_train, X_test, y_train, y_test = train_test_split(X_train, y_train, test_size=0.1

    X_train_ind = X_train['SK_ID_CURR']
    X_test_ind = X_test['SK_ID_CURR']

    X_valid.drop('SK_ID_CURR', inplace = True , axis = 1)
    X_train.drop('SK_ID_CURR', inplace = True , axis = 1)
    X_test.drop('SK_ID_CURR', inplace = True , axis = 1)

    return X_train, y_train, X_valid, y_valid, X_test, y_test, X_train_ind, X_val_ind,

```

In [67]:

```

def transform_train(X_train, X_valid, X_test, X_train_ind, X_val_ind, X_test_ind):
    num_attribs = ['AMT_INCOME_TOTAL', 'AMT_CREDIT', 'DAYS_EMPLOYED', 'DAYS_BIRTH', 'EXT_
                  'EXT_SOURCE_2', 'EXT_SOURCE_3']

    cat_attribs = ['CODE_GENDER', 'FLAG_OWN_REALTY', 'FLAG_OWN_CAR', 'NAME_CONTRACT_TYPE'
                  'NAME_EDUCATION_TYPE', 'OCCUPATION_TYPE', 'NAME_INCOME_TYPE']

    num_pipeline = Pipeline([
        ('selector', DataFrameSelector(num_attribs)),
        ('imputer', SimpleImputer(strategy='mean')),
        ('std_scaler', StandardScaler()),
    ])

    cat_pipeline = Pipeline([
        ('selector', DataFrameSelector(cat_attribs)),

```

```

        ('imputer', SimpleImputer(strategy='constant')),
        ('ohe', OneHotEncoder(sparse=False, handle_unknown="ignore"))
    ])

    data_prep_pipeline = FeatureUnion(transformer_list=[
        ("num_pipeline", num_pipeline),
        ("cat_pipeline", cat_pipeline),
    ])

    full_pipeline = Pipeline([("preparation", data_prep_pipeline)])

    X_train = pd.DataFrame(full_pipeline.fit_transform(X_train), index = X_train_ind)
    X_valid = pd.DataFrame(full_pipeline.transform(X_valid), index = X_val_ind)
    X_test = pd.DataFrame(full_pipeline.transform(X_test), index= X_test_ind)

    return full_pipeline, X_train, X_valid, X_test

```

## Baseline Model

In [48]:

```

def BaselineApproach(fitted_models, expLog, X_train, X_valid, X_test, y_train, y_test, y_val):
    models = [LogisticRegression(n_jobs = -1, solver = 'lbfgs'),
               RandomForestClassifier(n_jobs = -1, n_estimators = 100)]

    models_name = ['Logistic', 'RandomForest']
    ctr = 0
    p_value = 0
    for model in models:
        cv = ShuffleSplit(n_splits=30, test_size=0.3, random_state=0)
        np.random.seed(42)
        model.fit(X_train, y_train)
        exp_name = f"Model_{models_name[ctr]}_features_{X_train.shape[1]}"
        fitted_models.append(model)
        if ctr == 0:
            logit_scores = cross_val_score(model, X_train, y_train, cv=cv)
        else:
            best_train_scores = cross_val_score(model, X_train, y_train, cv=cv)
            (t_stat, p_value) = stats.ttest_rel(logit_scores, best_train_scores)

        expLog.loc[len(expLog)] = [f"{exp_name}"] + list(np.round(
            [accuracy_score(y_train, model.predict(X_train)),
             accuracy_score(y_valid, model.predict(X_valid)),
             accuracy_score(y_test, model.predict(X_test)),
             roc_auc_score(y_train, model.predict_proba(X_train)[: , 1]),
             roc_auc_score(y_valid, model.predict_proba(X_valid)[: , 1]),
             roc_auc_score(y_test, model.predict_proba(X_test)[: , 1]), p_value],
            4))

        ctr += 1
    return logit_scores

```

## Merge data from Previous Application file

In [62]:

```

def merge_df(X_train, X_valid, X_test, df_pv_app):
    X_train = X_train.merge(df_pv_app, how = 'left', on = 'SK_ID_CURR')
    X_valid = X_valid.merge(df_pv_app, how = 'left', on = 'SK_ID_CURR')
    X_test = X_test.merge(df_pv_app, how = 'left', on = 'SK_ID_CURR')

```

```

X_train.fillna(0, inplace = True)
X_valid.fillna(0, inplace = True)
X_test.fillna(0, inplace = True)

return X_train, X_valid, X_test

```

## Baseline model with Previous Application

```

In [51]: def BaselinPlusPvApp(logit_scores,fitted_models,expLog, X_train,X_valid, X_test,y_train
models = [LogisticRegression(n_jobs = -1, solver = 'lbfgs' )
          ,RandomForestClassifier(n_jobs = -1,n_estimators = 100)]
models_name = ['Logistic_Prev_app','RandomForest_prev_app']
ctr = 0
p_value = 0
for model in models:
    cv = ShuffleSplit(n_splits=30, test_size=0.3, random_state=0)
    np.random.seed(42)
    model.fit(X_train, y_train)
    exp_name = f"Model_{models_name[ctr]}_features_{X_train.shape[1]}"
    fitted_models.append(model)
    best_train_scores = cross_val_score(model, X_train, y_train, cv=cv)
    (t_stat, p_value) = stats.ttest_rel(logit_scores, best_train_scores)

    expLog.loc[len(expLog)] = [f"{exp_name}"] + list(np.round(
        [accuracy_score(y_train, model.predict(X_train)),
        accuracy_score(y_valid, model.predict(X_valid)),
        accuracy_score(y_test, model.predict(X_test)),
        roc_auc_score(y_train, model.predict_proba(X_train)[: , 1]),
        roc_auc_score(y_valid, model.predict_proba(X_valid)[: , 1]),
        roc_auc_score(y_test, model.predict_proba(X_test)[: , 1]),p_value],
        4))

    ctr += 1

```

## Performance Metrics

```

In [60]: fitted_models = []
expLog = pd.DataFrame(columns=["exp_name",
                              "Train Acc",
                              "Valid Acc",
                              "Test Acc",
                              "Train AUC",
                              "Valid AUC",
                              "Test AUC",
                              "P_Value"
                              ])

#Loading & Transforming Train
X_train, y_train, X_valid, y_valid, X_test, y_test, X_train_ind, X_val_ind, X_test_ind


full_pipeline, X_train, X_valid, X_test = transform_train(X_train, X_valid, X_test, X_t

#Baseline Model
logit_scores_baseline = BaselineApproach(fitted_models,expLog, X_train,X_valid, X_test,
expLog

```

Out[60]:

	exp_name	Train Acc	Valid Acc	Test Acc	Train AUC	Valid AUC	Test AUC	P_Value
0	Model_Logistic_features_48	0.9191	0.9192	0.9196	0.7357	0.7407	0.7450	0.0000
1	Model_RandomForest_features_48	0.9999	0.9194	0.9194	1.0000	0.7100	0.7169	0.1024



In [72]:

```
df_pv_app = previous_app()  
X_train, X_valid, X_test = merge_df(X_train, X_valid, X_test, df_pv_app)  
BaselinPlusPvApp(logit_scores_baseline,fitted_models,explog, X_train,X_valid, X_test,y_expLog
```

Out[72]:

	exp_name	Train Acc	Valid Acc	Test Acc	Train AUC	Valid AUC	Test AUC	P_Value
0	Model_Logistic_features_48	0.9191	0.9192	0.9196	0.7357	0.7407	0.7450	0.0000
1	Model_RandomForest_features_48	0.9999	0.9194	0.9194	1.0000	0.7100	0.7169	0.1024
2	Model_Logistic_Prev_app_features_269	0.9193	0.9194	0.9196	0.7567	0.7584	0.7601	0.0068
3	Model_RandomForest_prev_app_features_269	1.0000	0.9194	0.9195	1.0000	0.7168	0.7197	0.0000

In [70]:

```
datasets['previous_application'].describe()
```

Out[70]:

	SK_ID_CURR	AMT_ANNUITY	AMT_APPLICATION	AMT_CREDIT	AMT_DOWN_PAYMENT	AMT_G
count	1.670214e+06	1.297979e+06	1.670214e+06	1.670213e+06	7.743700e+05	1
mean	2.783572e+05	1.595512e+04	1.752339e+05	1.961140e+05	6.697402e+03	2
std	1.028148e+05	1.478214e+04	2.927798e+05	3.185746e+05	2.092150e+04	3
min	1.000010e+05	0.000000e+00	0.000000e+00	0.000000e+00	-9.000000e-01	0
25%	1.893290e+05	6.321780e+03	1.872000e+04	2.416050e+04	0.000000e+00	5
50%	2.787145e+05	1.125000e+04	7.104600e+04	8.054100e+04	1.638000e+03	1
75%	3.675140e+05	2.065842e+04	1.803600e+05	2.164185e+05	7.740000e+03	2



# Project Description

## Back ground Home Credit Group

Many people struggle to get loans due to insufficient or non-existent credit histories. And, unfortunately, this population is often taken advantage of by untrustworthy lenders.

### Home Credit Group

Home Credit strives to broaden financial inclusion for the unbanked population by providing a positive and safe borrowing experience. In order to make sure this underserved population has a positive loan experience, Home Credit makes use of a variety of alternative data--including telco and transactional information--to predict their clients' repayment abilities.

While Home Credit is currently using various statistical and machine learning methods to make these predictions, they're challenging Kagglers to help them unlock the full potential of their data. Doing so will ensure that clients capable of repayment are not rejected and that loans are given with a principal, maturity, and repayment calendar that will empower their clients to be successful.

## Background on the dataset

Home Credit is a non-banking financial institution, founded in 1997 in the Czech Republic.

The company operates in 14 countries (including United States, Russia, Kazakhstan, Belarus, China, India) and focuses on lending primarily to people with little or no credit history which will either not obtain loans or became victims of untrustworthy lenders.

Home Credit group has over 29 million customers, total assests of 21 billions Euro, over 160 millions loans, with the majority in Asia and and almost half of them in China (as of 19-05-2018).

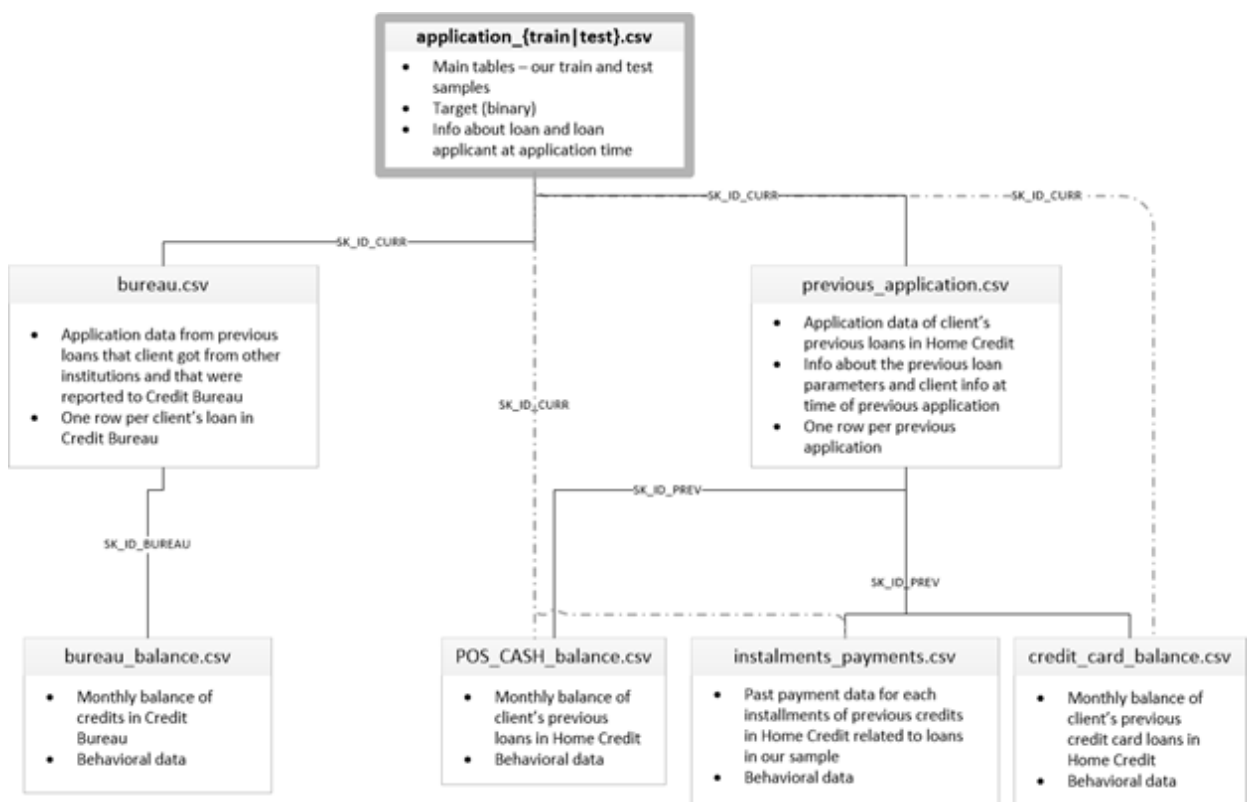
While Home Credit is currently using various statistical and machine learning methods to make these predictions, they're challenging Kagglers to help them unlock the full potential of their data. Doing so will ensure that clients capable of repayment are not rejected and that loans are given with a principal, maturity, and repayment calendar that will empower their clients to be successful.

## Data files overview

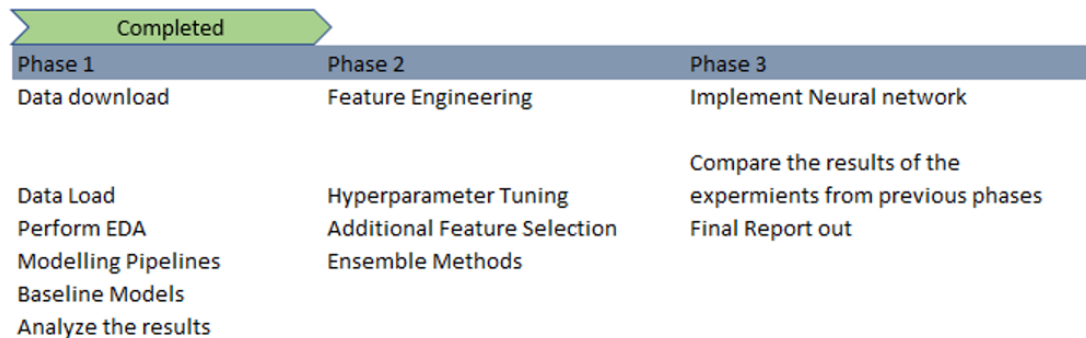
There are 7 different sources of data:

- **application\_train/application\_test:** the main training and testing data with information about each loan application at Home Credit. Every loan has its own row and is identified by the feature SK\_ID\_CURR. The training application data comes with the TARGET indicating **0: the loan was repaid** or **1: the loan was not repaid**. The target variable defines if the client had payment difficulties meaning he/she had late payment more than X days on at least one of the first Y installments of the loan. Such case is marked as 1 while other all other cases as 0.

- **bureau:** data concerning client's previous credits from other financial institutions. Each previous credit has its own row in bureau, but one loan in the application data can have multiple previous credits.
- **bureau\_balance:** monthly data about the previous credits in bureau. Each row is one month of a previous credit, and a single previous credit can have multiple rows, one for each month of the credit length.
- **previous\_application:** previous applications for loans at Home Credit of clients who have loans in the application data. Each current loan in the application data can have multiple previous loans. Each previous application has one row and is identified by the feature SK\_ID\_PREV.
- **POS\_CASH\_BALANCE:** monthly data about previous point of sale or cash loans clients have had with Home Credit. Each row is one month of a previous point of sale or cash loan, and a single previous loan can have many rows.
- **credit\_card\_balance:** monthly data about previous credit cards clients have had with Home Credit. Each row is one month of a credit card balance, and a single credit card can have many rows.
- **installments\_payment:** payment history for previous loans at Home Credit. There is one row for every made payment and one row for every missed payment.



## Project Phases and Activities



## Feature Engineering and transformers

Feature engineering is a process that can include both feature creation (adding new features to existing data) and feature selection (identifying only the most significant features or using other dimensionality reduction techniques). We may utilize a variety of ways to both build and pick features. Applied group\_by and grouped the data, taking SK\_ID\_CURR as a key, and aggregated using aggerate functions and using some features from the previous\_application data set.

We have performed correlation analysis in section 4.4 and picked up the highly correlated features for the experiments in this phase.

We have used the following Features

```
Selected features = ['SK_ID_CURR','AMT_INCOME_TOTAL',
'AMT_CREDIT','DAYS_EMPLOYED','DAYS_BIRTH','EXT_SOURCE_1',
'EXT_SOURCE_2','EXT_SOURCE_3','CODE_GENDER', 'FLAG_OWN_REALTY','FLAG_OWN_CAR',
'NAME_CONTRACT_TYPE','NAME_EDUCATION_TYPE','OCCUPATION_TYPE','NAME_INCOME_TYPE']
```

```
num_attribs = ['AMT_INCOME_TOTAL',
'AMT_CREDIT','DAYS_EMPLOYED','DAYS_BIRTH','EXT_SOURCE_1', 'EXT_SOURCE_2','EXT_SOURCE_3']
```

```
cat_attribs = ['CODE_GENDER', 'FLAG_OWN_REALTY','FLAG_OWN_CAR','NAME_CONTRACT_TYPE',
'NAME_EDUCATION_TYPE','OCCUPATION_TYPE','NAME_INCOME_TYPE']
```

## Pipelines

Build Numeric Pipeline

Identify the numeric features we wish to consider.

Create a pipeline for the numeric features.

Use DataFrameSelector() to select the desired numeric features



Use `SimpleImputer()` for missing data: there are quite a few missing values in the dataset. Missing values will be imputed using the feature mean.

Use `StandardScaler()` to standardize the data

The numeric pipeline will look like as follows:

```
num_pipeline = Pipeline([
    ('selector', DataFrameSelector(num_attribs)),
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler()),
])
```

Build Categorical Pipeline:

Identify the categorical features we wish to consider.

Identify the range of expected values for the categorical features.

Create a pipeline for the categorical features.

Use `SimpleImputer()` for missing data: there are quite a few missing values in the dataset. Missing values will be imputed using constant values.

Then use One Hot Encoding

The numeric pipeline will look like as follows:

```
cat_pipeline = Pipeline([
    ('selector', DataFrameSelector(cat_attribs)),
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('ohe', OneHotEncoder(sparse=False, handle_unknown="ignore"))
])
```

Union numeric pipeline and categorical pipeline:

The codes will looks like as follows:

```
data_prep_pipeline= FeatureUnion(transformer_list=[
    ("num_pipeline", num_pipeline),
    ("cat_pipeline", cat_pipeline),
])
```

```
full_pipeline = Pipeline([("preparation", data_prep_pipeline)])
```

# Experimental results

	exp_name	Train Acc	Valid Acc	Test Acc	Train AUC	Valid AUC	Test AUC	P_Value
0	Model_Logistic_features_48	0.9191	0.9192	0.9196	0.7357	0.7407	0.7450	0.0000
1	Model_RandomForest_features_48	0.9999	0.9194	0.9194	1.0000	0.7100	0.7169	0.1024
2	Model_Logistic_Prev_app_features_269	0.9193	0.9194	0.9196	0.7567	0.7584	0.7601	0.0068
3	Model_RandomForest_prev_app_features_269	1.0000	0.9194	0.9195	1.0000	0.7168	0.7197	0.0000

Initially, we got the test AUC of 0.74 and 0.71 with Logistic Regression and Random Forest, respectively, and after merging with the previous application\_data we could see a small increase in AUC in the Logistic Regression, but it had no effect in the Random Forest Model.

## Discussion

We have conducted 4 different experiments in this Phase using Logistic regression and Random Forrest algorithms. We have first conducted the experiment with teh data only from Application Train and Application Test data set with out including supporting datasets. We have observed that we got 74% AUC with logistic regression and 71% with Random forrest. We have then added additional features from Previous application data and observed that AUC for logist regression increased to 76% and haven't seen any improvement with Random forrest.

## Conclusion

By adding features from previous application data set , we have seen that accuracy improved from 74% to 76% for our baseline Logistic regression model. We haven't seen the improvement with the Random forest model. This shows the importance of doing EDA on other supporting datasets to find out the right features and use that in our models.

We followed a machine learning project's general outline:

- Recognize the issue and the data.
- Cleaning and formatting of data (this was mostly done for us)
- Examine exploratory data
- Model to start with: Baseline model
- Improved model
- Interpretation of the model (just a little)

Our focus in the coming phases will be on

Exploring new features


Hyperparameter tuning

Implement/explore other models like Gradient Boost and AdaBoost

# Kaggle Submission

Home Credit Default Risk

Can you predict how capable each applicant is of repaying a loan?

 Home Credit Group · 7176 teams · 3 years ago

\$70,000

Prize Money

Overview

Data

Code

Discussion

Leaderboard

Rules

Team

My Submissions

Late Submission

...

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
submission.csv	a few seconds ago	1 seconds	1 seconds	0.65311

Complete

[Jump to your position on the leaderboard](#)

You may select up to 2 submissions to be used to count towards your final leaderboard score. If 2 submissions are not selected, they will be automatically chosen based on your best submission scores on the public leaderboard. In the event that automatic selection is not suitable, manual selection instructions will be provided in the competition rules or by official forum announcement.

Your final score may not be based on the same exact subset of data as the public leaderboard, but rather a different private data subset of your full submission — your public score is only a rough indication of what your final score is.

You should thus choose submissions that will most likely be best overall, and not necessarily on the public subset.

In [ ]: